

Homework_6

Daria Barbour-Brown | Bailey Ho | Warren Kennedy

2023-05-10

Conceptual Problem

Question 1

When would you want to use ridge regression instead of a standard linear regression?

We should explore alternative methods to enhance both the accuracy and interpretability of our model predictions. One such method is ridge regression, which is a refined version of linear regression specifically designed for scenarios where the predictor variables exhibit high correlation. In situations where multicollinearity exists, traditional linear regression often yields unstable and unreliable coefficient estimates. By incorporating regularization, ridge regression effectively addresses this challenge by shrinking the coefficient estimates, reducing their variance, and stabilizing the model.

Question 2

When would you not want to use ridge regression?

Ridge regression uses regularization to shrink the coefficient estimates towards zero, although they never actually reach zero. This regularization can compromise interpretability, particularly if the aim is to interpret the significance of individual coefficients in your analysis. In such instances, alternative approaches like lasso regression may be more suitable. Lasso regression selectively identifies the most influential variables, simplifying the model and facilitating interpretation by emphasizing the most important predictors.

Application Question

Question 3

Part A

```
library(ISLR2)
data <- Hitters
data <- subset(data, select = -c(League, Division, NewLeague))
data <- data[complete.cases(data), ]
dim(data)
```

```
## [1] 263 17
```

Part B

When a variable's coefficient is reduced to zero by LASSO regression, it signifies that the variable does not possess any substantial influence on the model's prediction or that its impact is negligible when compared to other variables.

```
set.seed(1)
train <- sample(263, 263*.8)
```

Part C

```
lm.fit <- lm(Salary ~. , data = data, subset = train)
summary(lm.fit)
```

```
##
## Call:
## lm(formula = Salary ~ ., data = data, subset = train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -788.74 -177.36  -34.21  120.01 1912.33
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 177.00090   92.90453   1.905  0.0582 .
## AtBat       -1.74153    0.71980  -2.419  0.0165 *
## Hits         5.10323    2.85500   1.787  0.0754 .
## HmRun        -1.21466    6.91295  -0.176  0.8607
## Runs        -2.13662    3.42167  -0.624  0.5331
## RBI          3.08012    2.94988   1.044  0.2977
## Walks        4.02286    2.13287   1.886  0.0608 .
## Years       -15.47918   13.66545  -1.133  0.2587
## CAtBat       -0.24281    0.16309  -1.489  0.1382
## CHits        0.89063    0.80962   1.100  0.2727
## CHmRun       -0.39810    1.86142  -0.214  0.8309
## CRuns        0.73578    0.88780   0.829  0.4083
## CRBI         0.50915    0.84437   0.603  0.5472
## CWalks       -0.34568    0.41629  -0.830  0.4073
## PutOuts      0.44746    0.09015   4.963 1.52e-06 ***
## Assists      0.39990    0.23408   1.708  0.0892 .
## Errors      -4.67595    4.69122  -0.997  0.3201
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 315 on 193 degrees of freedom
## Multiple R-squared:  0.5259, Adjusted R-squared:  0.4866
## F-statistic: 13.38 on 16 and 193 DF,  p-value: < 2.2e-16
```

Part D

```
RMSE = sqrt(mean((data$Salary[-train] - predict(lm.fit, data[-train,]))^2))
RMSE
```

```
## [1] 395.7729
```

Part E

We anticipate that the Residual Standard Error (RSE) will be smaller, mainly because the Root Mean Squared Error (RMSE) is calculated by dividing the RSS by a larger number. Additionally, it is important to note that the RSS in Part C is derived from our training data, while the RSS in Part D pertains to the testing data. In general, we typically observe that the training RSS tends to be lower than the testing RSS. This occurrence is due to the fact that our model is specifically trained on the training data, which it analyzes during the training process. By doing so, it attempts to find the underlying patterns and relationships within that data. However, when the model encounters testing data with different patterns or relationships, it may struggle to capture the relationships that are unfamiliar to our model. Consequently, this disparity in data distribution can result in a comparatively higher testing RSS when contrasted with the training RSS.

Question 4

Part A

```
library(glmnet)
```

```
## Loading required package: Matrix
```


```
## Loaded glmnet 4.1-7
```

```
x <- model.matrix(Salary ~ ., data)[, -1]
y <- data$Salary

set.seed(1)
lasso.fit <- cv.glmnet(x[train, ], y[train], alpha = 1)
lasso.fit
```

```
##
## Call:  cv.glmnet(x = x[train, ], y = y[train], alpha = 1)
##
## Measure: Mean-Squared Error
##
##      Lambda Index Measure      SE Nonzero
## min    2.76    49 117059 25231         13
## 1se   94.61    11 138884 28841          5
```

```
plot(lasso.fit)
```



```
bestlam.lss <- lasso.fit$lambda.min
bestlam.lss
```

```
## [1] 2.757799
```

Part B

```
coef(lasso.fit)
```

```
## 17 x 1 sparse Matrix of class "dgCMatrix"
##                               s1
## (Intercept) 223.5746052
## AtBat      .
## Hits       0.5396439
## HmRun      .
## Runs       .
## RBI        .
## Walks      1.2242333
## Years      .
## CAtBat     .
## CHits      0.1174027
## CHmRun     .
## CRuns      .
```

```
## CRBI          0.2219827
## CWalks         .
## PutOuts       0.1721853
## Assists        .
## Errors         .
```

Part C

```
set.seed(1)
test <- (-train)
y.test <- y[test]

lasso.pred <- predict(lasso.fit, s = bestlam.lss, newx = x[test, ])
mean((lasso.pred - y.test)^2)
```

```
## [1] 155073.8
```


Question 5

Part A

```
ridge.fit <- cv.glmnet(x[train, ], y[train], alpha = 0)
ridge.fit
```

```
##
## Call:  cv.glmnet(x = x[train, ], y = y[train], alpha = 0)
##
## Measure: Mean-Squared Error
##
##      Lambda Index Measure      SE Nonzero
## min      24   100 116566 26201        16
## 1se     3322   47 141365 29937        16
```

```
plot(ridge.fit)
```



```
bestlam.rr <- ridge.fit$lambda.min
bestlam.rr
```

```
## [1] 23.98593
```

Part B

The intercept in ridge regression exhibits a greater magnitude compared to our linear estimates. Additionally, the coefficient estimates in ridge regression are, in many cases, closer to zero when compared to the estimates in our linear model.

```
coef(ridge.fit)
```

```
## 17 x 1 sparse Matrix of class "dgCMatrix"
##              s1
## (Intercept) 250.383896676
## AtBat      0.069678812
## Hits       0.286890025
## HmRun      0.906816756
## Runs       0.454298095
## RBI        0.476771168
## Walks      0.672986779
## Years      2.070773930
## CAtBat     0.007352379
```

```
## CHits      0.029216322
## CHmRun     0.172419128
## CRuns      0.055475339
## CRBI       0.055900249
## CWalks     0.058234872
## PutOuts    0.060078683
## Assists    0.010632032
## Errors     -0.152639646
```

```
coef(lm.fit)
```

```
## (Intercept)      AtBat      Hits      HmRun      Runs      RBI
## 177.0008998 -1.7415256  5.1032332 -1.2146591 -2.1366221  3.0801204
##      Walks      Years    CAtBat    CHits    CHmRun    CRuns
##   4.0228564 -15.4791836 -0.2428135  0.8906328 -0.3980951  0.7357843
##      CRBI      CWalks    PutOuts    Assists    Errors
##   0.5091505 -0.3456766  0.4474601  0.3999002 -4.6759469
```

Part C

```
ridge.pred <- predict(ridge.fit, s = bestlam.rr, newx = x[test, ])
ridge.pred
```

```
##              s1
## -Andre Dawson    949.85837
## -Andres Galarra  549.92155
## -Al Newman       106.43309
## -Andres Thomas   154.58464
## -Alex Trevino    315.09079
## -Barry Bonds     342.20301
## -Bill Buckner    1556.54955
## -Carlton Fisk    925.32168
## -Chris Speier    583.55044
## -Doug DeCinces   802.33537
## -Darrell Evans   1351.17563
## -Dan Gladden     371.68972
## -Dave Henderson  425.56860
## -Dale Murphy     933.26812
## -Don Slaught     473.66637
## -Eddie Milner    422.00707
## -Glenn Braggs    94.47897
## -George Brett    1151.53968
## -George Hendrick 895.41437
## -Gary Redus      345.50554
## -Gary Ward       658.28299
## -Howard Johnson  198.11654
## -Jose Cruz       1088.49440
## -Jeffrey Leonard 457.46062
## -Jerry Mumphy    619.66122
## -Jim Rice        1445.93949
## -Joel Skinner    207.40256
```

```
## -Kevin Bass      548.28906
## -Ken Griffey     859.06851
## -Ken Phelps      323.10433
## -Len Dykstra     442.60248
## -Lee Lacy        614.61700
## -Larry Sheets    238.44927
## -Mike Kingery    139.19749
## -Mike Marshall   346.30605
## -Ozzie Virgil    541.43760
## -Phil Bradley    583.72239
## -Paul Molitor    638.33103
## -Pete Rose       1740.04439
## -Pat Tabler      762.90128
## -Ron Hassey      519.72387
## -Rickey Henderson 884.50191
## -Ray Knight      623.68819
## -Rick Schu       140.94392
## -Steve Balboni   810.46817
## -Steve Garvey    1609.18080
## -Steve Jeltz     305.40352
## -Tony Gwynn      690.13904
## -Ted Simmons     1001.61748
## -Vince Coleman   359.85531
## -Wally Joyner    931.92586
## -Willie Randolph 923.46115
## -Wayne Tolleson  328.17524
```

```
mean((ridge.pred - y.test)^2)
```

```
## [1] 150137.8
```

Question 6

Which model would you recommend using if the General Manager of a baseball team is interested knowing which variables are most important for predicting a players salary?

In this scenario, our preference would be to utilize the LASSO estimates when reporting back to the General Manager. Although both LASSO and ridge regression can be employed for predictive modeling, LASSO's advantage lies in its ability to exclude less significant predictors, whereas ridge regression retains all predictors, thereby reducing model interpretability. Given the General Manager's primary interest in the most important variables, we opt for LASSO as it enables subset selection, emphasizing the variables of greatest relevance.