

Homework 8

Daria Barbour-Brown | Bailey Ho | Warren Kennedy

2023-05-28

Conceptual Problem

Question 1

Explain, in your own words, how a random forest makes predictions.

A random forest uses trees as building blocks to construct more powerful prediction models. This method is an adaptation of Bagging, i.e bootstrap aggregation, where we generate multiple bootstrapped training datasets, train a model on each set, and average all the predictions to obtain our bagging estimator. Random forests provide an improvement over bagged trees by randomizing the sample of predictors that can be used for splitting trees at each split of every bagged tree. This adaption decorrelates the bagged trees from each other increasing the variation accounted for by each tree, thereby making the average of the resulting trees less variable and more reliable. Our final random forest bagging estimator will be used for prediction.

Application Problems

Question 2

Using a linear model, predict bikers with the other variables in Bikeshare (excluding casual and registered). Make sure to use the variables in their appropriate form (i.e. numeric or factors as relevant to the context of the problem.) Perform variable selection as needed to get a strong predictive model. Use training/test splits of your choice and report a RMSE that describes your model's performance.

```
library(ISLR2)
library(MASS)
```

```
##
## Attaching package: 'MASS'

## The following object is masked from 'package:ISLR2':
##
## Boston
```

```
data <- Bikeshare

# convert Data
data$season <- as.factor(data$season)
data$hr <- as.numeric(data$hr)
```

```

data$holiday <- as.factor(data$holiday)
data$weekday <- as.factor(data$weekday)
data$workingday <- as.factor(data$workingday)

# Standardize the Data
numeric_columns <- sapply(data, is.numeric)
for (col in names(df)[numeric_columns]) {
  mean_val <- mean(df[[col]])
  sd_val <- sd(df[[col]])
  df[[col]] <- (df[[col]] - mean_val) / sd_val
}

#Train Linear Model
set.seed(1)
train <- sample(nrow(data), 0.8 * nrow(data))
train_set <- data[train, ]
test_set <- data[-train, ]
bikes.lm <- lm(bikers ~ . - casual - registered, data = train_set)
summary(bikes.lm)

```

```

##
## Call:
## lm(formula = bikers ~ . - casual - registered, data = train_set)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -246.80  -70.67  -20.20   44.28  430.19
##
## Coefficients: (1 not defined because of singularities)
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    10.0410     9.1081   1.102 0.270318
## season2        16.9472     7.8435   2.161 0.030755 *
## season3        29.1898     9.1912   3.176 0.001501 **
## season4        48.8306     7.7974   6.262 4.02e-10 ***
## mnthFeb        -2.8144     7.6531  -0.368 0.713073
## mnthMarch      -1.6140    10.8125  -0.149 0.881345
## mnthApril       11.1354    16.4281   0.678 0.497903
## mnthMay        43.1260    20.1609   2.139 0.032463 *
## mnthJune       11.4570    23.9704   0.478 0.632691
## mnthJuly      -18.4009    28.4856  -0.646 0.518318
## mnthAug         7.2090    32.2882   0.223 0.823331
## mnthSept       54.4690    35.9716   1.514 0.130016
## mnthOct        70.4675    40.2666   1.750 0.080158 .
## mnthNov        66.8146    44.4811   1.502 0.133120
## mnthDec        81.2873    48.3852   1.680 0.093002 .
## day           -0.2372     0.1430  -1.659 0.097109 .
## hr              5.6460     0.1925  29.336 < 2e-16 ***
## holiday1      -22.4822     7.7981  -2.883 0.003951 **
## weekday1        1.3630     4.7846   0.285 0.775757
## weekday2       -0.3855     4.7092  -0.082 0.934762
## weekday3       -1.2618     4.7551  -0.265 0.790749
## weekday4       -9.0617     4.7160  -1.922 0.054710 .
## weekday5        5.2910     4.6940   1.127 0.259705

```

```
## weekday6          5.3833      4.6597      1.155 0.248015
## workingday1         NA          NA          NA          NA
## weathersitcloudy/misty 7.3621      3.0810      2.390 0.016897 *
## weathersitlight rain/snow -21.8856      5.0174     -4.362 1.31e-05 ***
## weathersitheavy rain/snow 9.9748     103.8335      0.096 0.923471
## temp              209.9153      59.0607      3.554 0.000382 ***
## atemp             109.6514      62.6714      1.750 0.080228 .
## hum               -160.5908      8.2685     -19.422 < 2e-16 ***
## windspeed         27.4018      11.6728      2.347 0.018929 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 103.6 on 6885 degrees of freedom
## Multiple R-squared:  0.4078, Adjusted R-squared:  0.4052
## F-statistic: 158.1 on 30 and 6885 DF, p-value: < 2.2e-16
```

```
lm.model <- stepAIC(bikes.lm, direction = "both")
```

```
## Start: AIC=64217.57
## bikers ~ (season + mnth + day + hr + holiday + weekday + workingday +
##     weathersit + temp + atemp + hum + windspeed + casual + registered) -
##     casual - registered
##
##
## Step: AIC=64217.57
## bikers ~ season + mnth + day + hr + holiday + weekday + weathersit +
##     temp + atemp + hum + windspeed
##
##           Df Sum of Sq    RSS   AIC
## <none>             73882885 64218
## - day           1     29544 73912429 64218
## - weekday       6    139058 74021944 64219
## - atemp         1     32850 73915735 64219
## - windspeed     1     59135 73942021 64221
## - holiday       1     89194 73972079 64224
## - temp          1    135560 74018445 64228
## - weathersit     3    361153 74244039 64245
## - season        3    422999 74305885 64251
## - mnth         11   1291592 75174478 64315
## - hum           1   4047910 77930795 64584
## - hr            1   9234864 83117749 65030
```

```
summary(lm.model)
```

```
##
## Call:
## lm(formula = bikers ~ season + mnth + day + hr + holiday + weekday +
##     weathersit + temp + atemp + hum + windspeed, data = train_set)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -246.80  -70.67  -20.20   44.28  430.19
##
```

```
## Coefficients:
##               Estimate Std. Error t value Pr(>|t|)
## (Intercept)    10.0410     9.1081   1.102 0.270318
## season2        16.9472     7.8435   2.161 0.030755 *
## season3        29.1898     9.1912   3.176 0.001501 **
## season4        48.8306     7.7974   6.262 4.02e-10 ***
## mnthFeb        -2.8144     7.6531  -0.368 0.713073
## mnthMarch      -1.6140    10.8125  -0.149 0.881345
## mnthApril       11.1354    16.4281   0.678 0.497903
## mnthMay         43.1260    20.1609   2.139 0.032463 *
## mnthJune        11.4570    23.9704   0.478 0.632691
## mnthJuly       -18.4009    28.4856  -0.646 0.518318
## mnthAug         7.2090    32.2882   0.223 0.823331
## mnthSept        54.4690    35.9716   1.514 0.130016
## mnthOct         70.4675    40.2666   1.750 0.080158 .
## mnthNov         66.8146    44.4811   1.502 0.133120
## mnthDec         81.2873    48.3852   1.680 0.093002 .
## day            -0.2372     0.1430  -1.659 0.097109 .
## hr              5.6460     0.1925  29.336 < 2e-16 ***
## holiday1       -22.4822     7.7981  -2.883 0.003951 **
## weekday1        1.3630     4.7846   0.285 0.775757
## weekday2       -0.3855     4.7092  -0.082 0.934762
## weekday3       -1.2618     4.7551  -0.265 0.790749
## weekday4       -9.0617     4.7160  -1.922 0.054710 .
## weekday5        5.2910     4.6940   1.127 0.259705
## weekday6        5.3833     4.6597   1.155 0.248015
## weathersitcloudy/misty 7.3621     3.0810   2.390 0.016897 *
## weathersitlight rain/snow -21.8856     5.0174  -4.362 1.31e-05 ***
## weathersitheavy rain/snow 9.9748    103.8335   0.096 0.923471
## temp          209.9153     59.0607   3.554 0.000382 ***
## atemp         109.6514     62.6714   1.750 0.080228 .
## hum          -160.5908     8.2685 -19.422 < 2e-16 ***
## windspeed      27.4018     11.6728   2.347 0.018929 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 103.6 on 6885 degrees of freedom
## Multiple R-squared:  0.4078, Adjusted R-squared:  0.4052
## F-statistic: 158.1 on 30 and 6885 DF, p-value: < 2.2e-16
```

```
#Test Model
predict_test <- predict(lm.model, newdata = test_set)
rmse <- sqrt(mean((test_set$bikers - predict_test)^2))
rmse
```

```
## [1] 101.5678
```

Question 3

Explain your approach in 1, including how you coded variables (as numeric or factor), variable selection procedure, and training/test splits you decided to use.

In question one we first changed the data frame to ensure that our data was being interpreted correctly. This involved converting hours to numeric, and switching “holiday, weekday, workingday, and season” to factor

variables with corresponding levels. Then we converted variables to factor if their levels were not continuously defined. For example, season was coded with numbers, but these numbers are only defined when they are integers, so we converted the variables to make sure there were no ill-defined outcomes. Our process was to train a linear model using our training data and then perform subset selection using both forward and backward step wise selection to decide which variables were most important. We ended up using the validation set approach to test our method, withholding 20 percent of our data in order to test our model's predictive power.

Question 4

Using a random forest, predict bikers with the other variables in Bikeshare (excluding casual and registered). Make sure to use the variables in their appropriate form (i.e. numeric or factors as relevant to the context of the problem.) Perform model tuning to get a strong predictive model. Use training/test/validation splits of your choice and report a RMSE that describes your model's performance.

```
# Train Random F0rest
library(caret)

## Loading required package: ggplot2

## Loading required package: lattice

library(randomForest)

## randomForest 4.7-1.1

## Type rfNews() to see new features/changes/bug fixes.

##
## Attaching package: 'randomForest'

## The following object is masked from 'package:ggplot2':
##
##      margin

set.seed(1)
train <- sample(nrow(data), 0.8 * nrow(data))
train_set <- data[train, ]
test_set <- data[-train, ]
bikes.rf <- randomForest(bikers ~ . - casual - registered, data = train_set, mtry = 6, importance = TRUE)

#Test Random Forest
test_predict <- predict(bikes.rf, newdata = test_set)
rmse <- sqrt(mean((test_set$bikers - test_predict)^2))
rmse

## [1] 37.18259
```

```

# Pruning/Cross Validation

### Define the control parameters for cross-validation
ctrl <- trainControl(method = "cv", number = 5, verboseIter = FALSE) # 5-fold cross-validation

### Train the random forest model using cross-validation
model.rf2 <- train(bikers ~ . - casual - registered, data = train_set, method = "rf", trControl = ctrl)

### Access the cross-validation results
cv_results <- model.rf2$results
cv_results

##      mtry      RMSE Rsquared      MAE  RMSESD RsquaredSD      MAESD
## 1      2 94.64802 0.5878814 71.11060 1.194627 0.023456148 0.4909656
## 2     16 36.43452 0.9288495 23.81025 2.067327 0.010690141 0.7116789
## 3     31 34.47204 0.9341856 21.68906 1.269823 0.007885316 0.5087632

#Test Random Forest
test_predict <- predict(model.rf2, newdata = test_set)
rmse2 <- sqrt(mean((test_set$bikers - test_predict)^2))
rmse2

## [1] 37.5201

```

Question 5

Explain your approach in 3, including how you coded variables (as numeric or factor), variable selection procedure, and training/test splits you decided to use.

For our random forest model, we retained the same data frame transformations that we had put in place in question 1. For this model, we first trained the random forest on the training data, setting our “mtry” to 6, meaning that we restricted our split decision to subsets of 6 features for each split in each tree. We then trained another random forest using 5-fold cross validation in order to ensure better tune our model in an attempt to improve our models predictive power. In training our random forest, we used essentially the same training and test set as was used in our linear model.

Question 6

Plot variable importances of your final random forest. Compare these with the linear regression output. Does the random forest find the same variables to be important as the linear regression?

```

importance(bikes.rf)

##           %IncMSE IncNodePurity
## season      25.11616      2001296.9
## mnth        25.39012      4022040.1
## day         44.14170      6329836.2
## hr          315.40463      67307635.9
## holiday      18.30377       320532.2
## weekday      53.57751      5323534.8
## workingday   62.23800      4237916.8

```

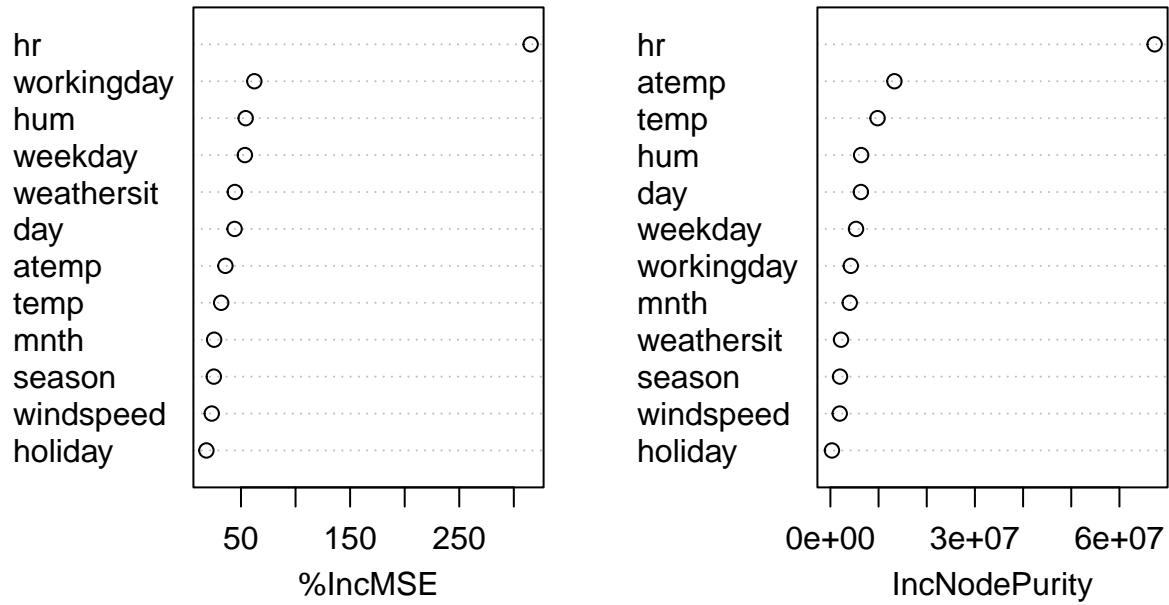
```
## weathersit 44.36976 2206503.3
## temp      31.88275 9792371.1
## atemp     35.71790 13294070.6
## hum       54.30114 6382761.2
## windspeed 23.21635 1938222.9
```

```
bikes.lm
```

```
##
## Call:
## lm(formula = bikers ~ . - casual - registered, data = train_set)
##
## Coefficients:
##              (Intercept)                season2
##              10.0410                16.9472
##              season3                season4
##              29.1898                48.8306
##              mnthFeb                mnthMarch
##              -2.8144                -1.6140
##              mnthApril                mnthMay
##              11.1354                43.1260
##              mnthJune                mnthJuly
##              11.4570                -18.4009
##              mnthAug                mnthSept
##              7.2090                54.4690
##              mnthOct                mnthNov
##              70.4675                66.8146
##              mnthDec                day
##              81.2873                -0.2372
##              hr                holiday1
##              5.6460                -22.4822
##              weekday1                weekday2
##              1.3630                -0.3855
##              weekday3                weekday4
##              -1.2618                -9.0617
##              weekday5                weekday6
##              5.2910                5.3833
##              workingday1    weathersitcloudy/misty
##              NA                7.3621
## weathersitlight rain/snow    weathersitheavy rain/snow
##              -21.8856                9.9748
##              temp                atemp
##              209.9153                109.6514
##              hum                windspeed
##              -160.5908                27.4018
```

```
varImpPlot(bikes.rf)
```

bikes.rf



Both linear regression and random variables find the “hr” variable(s) to be important, From there on, it is tough to distinguish which variables are most influential in both models considering because the coefficients for the linear model tend to be clustered around similar values, and in similar fashion, the measured importance of the variables in the random forest are nearly identical.